


Deep learning for medical imaging school

Hands-on session

Autoencoders

By
Pierre-Marc Jodoin, Thierry Judge



As you can see, there is
no code and no data

Welcome to a
brand new empty
session!

Click here

In the terminal, type

```
git clone https://github.com/vitalab/deep-learning-tutorials.git
```



This command allows to **download the code** from a public **gitHub repository**



File Edit View Run Kernel Git Tabs Settings Help

Filter files by name

/workspace/deep-learning-tutorials/


Name	Last Modified
src	5 minutes ago
tutorials	5 minutes ago
dev.txt	5 minutes ago
environment.yml	5 minutes ago
LICENSE	5 minutes ago
pyproject.toml	5 minutes ago
README.md	5 minutes ago
requirements.txt	4 minutes ago
setup.cfg	5 minutes ago
setup.py	5 minutes ago


Terminal 1 requirements.txt

```
jovyan@w-pierr-autoencoderhandson2-1fc514f792564bc383f7cbda623718zpdgk:~/workspace$ git clone https://github.com/vitalab/deep-learning-tutorials.git
Cloning into 'deep-learning-tutorials'...
remote: Enumerating objects: 280, done.
remote: Counting objects: 100% (198/198), done.
remote: Compressing objects: 100% (149/149), done.
remote: Total 280 (delta 134), reused 84 (delta 49), pack-reused 82
Receiving objects: 100% (280/280), 66.88 KiB | 5.14 MiB/s, done.
Resolving deltas: 100% (156/156), done.
jovyan@w-pierr-autoencoderhandson2-1fc514f792564bc383f7cbda623718zpdgk:~/workspace$
```

Great! The code is there!

Lets download the data





Download the following file on your computer

<https://drive.google.com/file/d/1H5pTOYjcSFR6B5GhA0sEPW0wgPVfBq8S/view?usp=sharing>



File Edit View Run Kernel Git Tabs Settings Help



Filter files by name


/ workspace / deep-learning-tutorials /

Name	Last Modified
src	6 minutes ago
tutorials	6 minutes ago
data.tar.gz	a year ago
dev.txt	6 minutes ago
environment.yml	6 minutes ago
LICENSE	6 minutes ago
pyproject.toml	6 minutes ago
README.md	6 minutes ago
requirements.txt	6 minutes ago
se	
se	

data files should
appear here

Terminal 1 requirements.txt

```
jovyan@pierr-autoencoderhandson2-1fc514f792564bc383f7cbda623718zpdgk:~/workspace$ git clone https://github.com/vitalab/deep-learning-tutorials.git
Cloning into 'deep-learning-tutorials'...
remote: Enumerating objects: 280, done.
remote: Counting objects: 100% (198/198), done.
remote: Compressing objects: 100% (149/149), done.
remote: Total 280 (delta 134), reused 84 (delta 49), pack-reused 82
Receiving objects: 100% (280/280), 66.88 KiB | 4.78 MiB/s, done.
Resolving deltas: 100% (156/156), done.
jovyan@pierr-autoencoderhandson2-1fc514f792564bc383f7cbda623718zpdgk:~/workspace$ cd deep-learning-tutorials/
jovyan@pierr-autoencoderhandson2-1fc514f792564bc383f7cbda623718zpdgk:~/workspace/deep-learning-tutorials$ wget --no-check-certificate http://info.usherbrooke.ca/pmjodoin/projects/data.tar.gz
--2022-04-19 10:03:10-- http://info.usherbrooke.ca/pmjodoin/projects/data.tar.gz
Resolving info.usherbrooke.ca (info.usherbrooke.ca)... 132.210.238.207
Connecting to info.usherbrooke.ca (info.usherbrooke.ca)|132.210.238.207|:80... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://info.usherbrooke.ca/pmjodoin/projects/data.tar.gz [following]
--2022-04-19 10:03:10-- https://info.usherbrooke.ca/pmjodoin/projects/data.tar.gz
Connecting to info.usherbrooke.ca (info.usherbrooke.ca)|132.210.238.207|:443... connected.
WARNING: cannot verify info.usherbrooke.ca's certificate, issued by 'CN=R3,0=Let's Encrypt,C=US':
Unable to locally verify the issuer's authority.
HTTP request sent, awaiting response... 200 OK
189853183 (181M) [application/x-gzip]
to: 'data.tar.gz'
100%[=====] 181.06M 56.8MB/s in 3.5s
--19 10:03:14 (51.6 MB/s) - 'data.tar.gz' saved [189853183/189853183]
pierr-autoencoderhandson2-1fc514f792564bc383f7cbda623718zpdgk:~/workspace/deep-learning-tutorials$
```

File Edit View Run Kernel Git Tabs Settings Help

Filter files by name	
/ workspace / deep-learning-tutorials /	
Name	Last Modified
data	a year ago
src	8 minutes ago
tutorials	8 minutes ago
data.tar.gz	a year ago
dev.txt	8 minutes ago
environment.yml	8 minutes ago
LICENSE	8 minutes ago
pyproject.toml	8 minutes ago
README.md	8 minutes ago
requirements.txt	8 minutes ago
setup.cfg	8 minutes ago
setup.py	8 minutes ago

Terminal 1 requirements.txt

```
jovyan@w-pierr-autoencoderhandson2-1fc514f792564bc383f7cbda623718zpdgk:~/workspace$ git clone https://github.com/vitalab/deep-learning-tutorials.git
Cloning into 'deep-learning-tutorials'...
remote: Enumerating objects: 280, done.
remote: Counting objects: 100% (198/198), done.
remote: Compressing objects: 100% (149/149), done.
remote: Total 280 (delta 134), reused 84 (delta 49), pack-reused 82
Receiving objects: 100% (280/280), 66
Resolving deltas: 100% (156/156), don
jovyan@w-pierr-autoencoderhandson2-1f
jovyan@w-pierr-autoencoderhandson2-1f
--2022-04-19 10:03:10-- http://info.
Resolving info.usherbrooke.ca (info.u
Connecting to info.usherbrooke.ca (in
HTTP request sent, awaiting response.
Location: https://info.usherbrooke.ca
--2022-04-19 10:03:10-- https://info
Connecting to info.usherbrooke.ca (in
WARNING: cannot verify info.usherbroo
Unable to locally verify the issuer
HTTP request sent, awaiting response.
Length: 189853183 (181M) [application
Saving to: 'data.tar.gz'
```

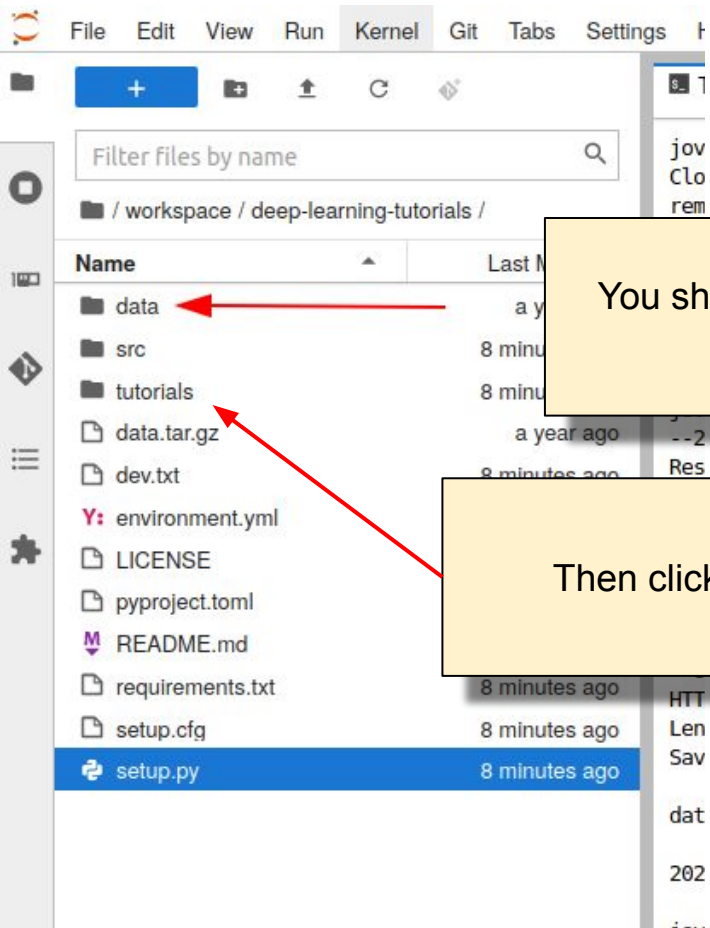
Then in the terminal, type

```
tar -xvzf data.tar.gz
```

data.tar.gz

2022-04-19 10:03:14 (51.6 MB/s) - 'data.tar.gz' saved [189853183/189853183]

```
jovyan@w-pierr-autoencoderhandson2-1fc514f792564bc383f7cbda623718zpdgk:~/workspace/deep-learning-tutorials$ tar -xvzf data.tar.gz
data/
data/acdc.h5
data/MNIST/
data/MNIST/raw/
data/MNIST/raw/train-labels-idx1-ubyte
data/MNIST/raw/train-images-idx3-ubyte.gz
data/MNIST/raw/train-labels-idx1-ubyte.gz
data/MNIST/raw/t10k-images-idx3-ubyte.gz
data/MNIST/raw/t10k-labels-idx1-ubyte
data/MNIST/raw/t10k-images-idx3-ubyte
data/MNIST/raw/train-images-idx3-ubyte
data/MNIST/raw/t10k-labels-idx1-ubyte.gz
data/MNIST/processed/
data/MNIST/processed/test.pt
data/MNIST/processed/training.pt
jovyan@w-pierr-autoencoderhandson2-1fc514f792564bc383f7cbda623718zpdgk:~/workspace/deep-learning-tutorials$
```



You should have the following
data folder

Then click on **tutorials**

File Edit View Run Kernel Git Tabs Settings Help

Filter files by name

/ ... / deep-learning-tutorials / tutorials /

Name	Last Modified
cardiac-mri-autoencoders.i...	9 minutes ago
mnist-autoencoders.ipynb	9 minutes ago

Terminal 1

requirements.txt

```
joyvan@w-pierr-autoencoderhandson2-1fc514f792564bc383f7cbda623718zpdgk:~/workspace$ git clone https://github.com/vit  
Cloning into 'deep-learning-tutorials'...
```

These are the two Jupyter notebooks for this tutorial.

Now you are good to go!

Start with **mnist** and then move on to **cardiac-mri**

```
Unabl  
HTTP request sent, awaiting response... 200 OK  
Length: 189853183 (181M) [application/x-gzip]  
Saving to: 'data.tar.gz'
```

Autoencoders'

Hands on

This hands-on has **two tutorials**.

The **second one** shows how to encode a **cardiac shapes** (ACDC dataset) with an autoencoder and a variational autoencoder.

Please see the last pages of this document for **more details on cardiac shapes**

```
mnist-autoencoders.ipynb × cardiac-mri-autoencoders.ipynb × latent_space.py × +
Filter files by name
/ tutorials /
Name Modified
cardiac-mri-autoe... 3 min. ago
mnist-autoencod... 2 min. ago

[14]: from src.visualization.latent_space import explore_latent_space

explore_latent_space(
    acdc_val,
    num_classes=4).float()[ :, :latent_space_size],
```


the 32D latent space




Decoded sample




Autoencoders (recap)




Annotated dataset: $D = \{(\vec{x}_1, t_1), (\vec{x}_2, t_2), \dots, (\vec{x}_N, t_N)\}$






Annotated dataset: $D = \{(\vec{x}_1, t_1), (\vec{x}_2, t_2), \dots, (\vec{x}_N, t_N)\}$




Annotated dataset: $D = \{(\vec{x}_1, t_1), (\vec{x}_2, t_2), \dots, (\vec{x}_N, t_N)\}$





GAN

(Latent space)




Loss?



Annotated dataset?

GAN

(Latent space)




Autoencoders



Autoencoders (once training is over)

(Latent space)




Summary


1. What are autoencoders and variational autoencoders?
2. How do they apply to MNIST (grayscale images)?
3. How do they apply to segmentation maps (ACDC cardiac labels)?


Goal : learn the latent representation of a set of data


How : by training a Neural Net to output its own ...
input!

Note : if you are familiar with AE and VAE, you may skip the next couple of slides.











Fully-Connected Layers

h : activation function




Very often...

No activation function at the output
of the encoder and the decoder



See **why?**


Sometimes **sigmoid** to predict pixel values between 0 and 1 or **ReLU** when the pixel values can be large but never negative.



The number of neurons

Decrease (or stay the same)


from a layer to another



The number of neurons


Increase (or stay the same)

from a layer to another





Very often...


The structure of the encoder is the **dual** of that of the decoder



Basic image-based autoencoder



$\vec{z} \in R^d$



Simple MNIST Autoencoder

```
class autoencoder(nn.Module):  
    def __init__(self):  
        super(autoencoder, self).__init__()  
        self.encoder = nn.Sequential(  
            nn.Linear(28 * 28, 128), nn.ReLU(True),  
            nn.Linear(128, 64), nn.ReLU(True),  
            nn.Linear(64, 12), nn.ReLU(True),  
            nn.Linear(12, 2))  
        self.decoder = nn.Sequential(  
            nn.Linear(2, 12), nn.ReLU(True),  
            nn.Linear(12, 64), nn.ReLU(True),  
            nn.Linear(64, 128), nn.ReLU(True),  
            nn.Linear(128, 28 * 28))  
  
    def forward(self, x):  
        z = self.encoder(x)  
        x_prime = self.decoder(z)  
        return x_prime
```

Latent space 2D



Simple MNIST Autoencoder


```
class autoencoder(nn.Module):  
    def __init__(self):  
        super(autoencoder, self).__init__()  
        self.encoder = nn.Sequential(  
            nn.Linear(28 * 28, 128), nn.ReLU(True),  
            nn.Linear(128, 64), nn.ReLU(True),  
            nn.Linear(64, 12), nn.ReLU(True),  
            nn.Linear(12, 2))  
        self.decoder = nn.Sequential(  
            nn.Linear(2, 12), nn.ReLU(True),  
            nn.Linear(12, 64), nn.ReLU(True),  
            nn.Linear(64, 128), nn.ReLU(True),  
            nn.Linear(128, 28 * 28))  
  
    def forward(self, x):  
        z = self.encoder(x)  
        x_prime = self.decoder(z)  
        return x_prime
```

symmetry




MNIST latent space (for 1000 images)


Each 2D point corresponds to an image




Conv layers





Basic autoencoder



Variational autoencoder





Variational autoencoder




The encoder outputs a **distribution** $p(\vec{z}|\vec{x})$
and not just a **vector** $\vec{z}$

Variational autoencoder




Random sample $\vec{z} \sim P(\vec{z}|\vec{x})$

Variational autoencoder




...and rebuild $\hat{x}$

$$p(\vec{z}|\vec{x}) \sim \textit{Gaussian}$$




$$p(\vec{z}|\vec{x}) \sim \text{Gaussian}$$




To simplify things, we assume Σ is **diagonal**

$$\Sigma = \begin{pmatrix} \sigma_1^2 & 0 & 0 & 0 \\ 0 & \sigma_1^2 & 0 & 0 \\ 0 & 0 & \ddots & 0 \\ 0 & 0 & 0 & \sigma_d^2 \end{pmatrix}$$


Variational autoencoder




Another way of seeing things...



Another way of seeing things...




ELBO loss : Evidence Lower Bound loss




$$Loss = \underbrace{(\vec{x} - \hat{x})^2}_{\text{Loss decoder}} + \lambda \underbrace{KL\left(N(\vec{z}; \vec{0}, \vec{1}), N(\vec{z}; \vec{\mu}, \vec{\Sigma}) \right)}_{\text{Loss encoder}}$$

Other loss (in case the output is binary)




$$\text{Loss} = \underbrace{\text{CrossEntropy}(\vec{x}, \hat{x})}_{\text{Loss decoder}} + \lambda \underbrace{\text{KL}\left(\text{N}(\vec{z}; \vec{0}, \vec{1}), \text{N}(\vec{z}; \vec{\mu}, \vec{\Sigma})\right)}_{\text{Loss encoder}}$$


Reparametrization trick instead of sampling



Basic autoencoder




Variational autoencoder




MNIST

Visualize the latent space (autoencoder)



MNIST

Visualize the latent space (variational autoencoder)



MNIST

Code is in the file:

`mnist-autoencoders.ipynb`


Automated Cardiac Diagnosis Challenge



150 exams (all from different patients) divided into 5 groups:

- dilated cardiomyopathy (DCM),
- hypertrophic cardiomyopathy (HCM),
- myocardial infarction (MINF),
- abnormal right ventricle (RV)
- patients without cardiac disease (NOR).

website: ***creatis.insa-lyon.fr/Challenge/acdc***



Overview
Contest
Participation
Databases
Evaluation
Code
Paper Submission
Contact

Overview

[General context](#) [Scientific interests](#) [Organizers](#)

The goal of this contest is two-fold:

- compare the performance of automatic methods on the segmentation of the left ventricular endocardium and epicardium as the right ventricular endocardium for both end diastolic and end systolic phase instances;
- compare the performance of automatic methods for the classification of the examinations in five classes (normal case, heart failure with infarction, dilated cardiomyopathy, hypertrophic cardiomyopathy, abnormal right ventricle).

This will be done using a common database of 3D cine-MRI images acquired from 150 patients (30 per pathology plus 30 healthy subjects) and the associated manual references based on the analysis of a clinical expert.

NEWS

4th of April 2017


The full training dataset (100 patients) is available

13th of March 2017

Challenge accepted by the MICCAI Satellite Committee


10th of January 2017

The online evaluation platform is operational



Cardiac anatomy


Normal Heart



Chambers relax and fill,
then contract and pump.


Cardiac anatomy

Normal Heart




Chambers relax and fill,
then contract and pump.

Cross-section : **short axis view**






Short axis cardiac MRI




Short axis segmentation map






Other examples:



Convolutional Variational autoencoder



ACDC

Code is in the file:

cardiac-mri-autoencoders.ipynb

Thank you